

## Lab 3:

### Virtual Realities

Due March 23 by the start of lecture.

#### Overview

In this lab, you will explore the dynamic binding mechanism of virtual functions in object-oriented languages. We will hand-build a vtable for a group of “classes” in C, and show how the compiler uses a vtable to select a derived function to call at run time based on the type of the value it is called on.

#### Function Pointers

C and C++ allow programmers to declare variables that point to *functions* rather than *values*. When such a pointer is dereferenced, the actual function that it points to is called, using arguments supplied by the programmer. C and C++ do type checking of function pointers to make sure the right number and types of arguments are provided at compile time. Example:

```
int max(int a, int b) { return a >= b ? a : b; } // return the larger integer

// in main():
int (*pFunc)(int, int);    // pFunc is a pointer to a function that takes 2 int arguments
                           // and returns int
pFunc = max;               // pFunc now points to the real function “max”
printf(“%d”, pFunc(5, 10)); // dereference pFunc, use its return value
```

In F#, we say that a function taking two integers and returning integer has type `int->int->int`. In C, we say it has type `int (*)(int, int)`. (As if we replaced the name of the function with `(*)`, removed the parameter names, and left everything else the same.)

Like all pointers, we do not necessarily know at compile time what a function pointer actually points to, so the actual function to be executed will not be known until run-time. We can use this with some trickery to implement function calls that execute different function bodies depending on run-time values.

#### Faking Objects and Polymorphism

We will implement this lab in C, which means we do not have access to objects... but we can fake it with structs. After all, a struct is a class without member functions, inheritance, or polymorphism... but it *can* store a pointer to a vtable, and through this we will implement polymorphism and dynamic dispatch.

Consider the following **C++ code** showing the polymorphism feature we want to emulate:

```
class Employee {
    int age;
public:
    int GetAge() { return age; }
    virtual void Speak() = 0;
    virtual double GetPay() = 0;
};

class HourlyEmployee : public Employee {
    double hourly_rate;
    double hours;
```

```

public:
    void Speak() { cout << "I work for " << hourly_rate << " dollars per hour :(";
};

class CommissionEmployee : public Employee {
    double sales_amount;
public:
    void Speak() { cout << "I make commission on " << sales_amount << " dollars in sales!";
};

// in main()
Employee *e = ... // suppose e points to either a CommissionEmployee or HourlyEmployee.
cout << e->GetPay(); // which function gets called? Who knows!

```

We can emulate this in C using:

1. A `struct Employee`, with two member fields: a **pointer** to vtable (as a `void**`); and an integer field `age`.
2. A `struct HourlyEmployee`, with four member fields: the same fields as `Employee`, and doubles `hourly_rate` and `hours`.
3. A `struct CommissionEmployee`, likewise, but with a double `sales_amount`.
4. Global functions to emulate each of the member functions of the objects:
  - (a) `Speak_Hourly` which takes an `Employee` pointer, casts it to a `HourlyEmployee` pointer, and prints the employee's message;
  - (b) `GetPay_Hourly` which also takes an `Employee` pointer and returns the employee's total pay (see below);
  - (c) `Construct_Hourly` which takes a `HourlyEmployee` pointer and initializes its fields to their default values, **most importantly** initializing the object's vtable pointer (more on this below).

and then following some tedious steps to create variables of our types and call the appropriate functions:

1. To make a `HourlyEmployee`, declare a `HourlyEmployee` variable (either on the stack or with `malloc`) and then pass it by pointer to `Construct_Hourly`.
2. To use "polymorphism" to point an `Employee` at a `HourlyEmployee`, declare an `Employee` pointer and initialize it by casting your `HourlyEmployee`'s address to an `Employee` pointer.
3. To use "dynamic dispatch", dereference the `Employee` pointer's vtable pointer, index the table to the appropriate method, cast that pointer to the correct function pointer type, and invoke the method by passing the `Employee` pointer and any other necessary parameters. Easy!

## Vtables

A **vtable** (short for **virtual table**) is a table of function pointers, with one entry in the table for each virtual function in a class (or its ancestors). Since the `Employee` class has two virtual functions, the vtable for any `Employee`-derived object will have two pointers in it for `Speak` and `GetCost`, plus additional entries for any more virtual functions introduced by the derived class. C++ handles the creation and use of vtables automatically; in this lab, we will simulate the work that a C++ compiler performs to transform virtual method calls into vtable lookups.

Suppose we have two functions: `void Speak_Hourly(struct Employee*)` and `double GetPay_Hourly(struct Employee*)`. We can build a "table" (really an array) of two pointers to functions as such:

```
void* Vtable_Hourly[2] = {Speak_Hourly, GetPay_Hourly};
```

The `void*` type in C lets us create a pointer to anything, but the compiler won't help us use such a pointer; we will need to cast it to something specific in order to use it. Suppose we have a `struct HourlyEmployee h` variable that we want to call `Speak_Hourly` on, but we don't want to use `Speak_Hourly` directly – instead, we want to use `Vtable_Hourly`. We note that the first entry in `Vtable_Hourly` is a pointer to `Speak_Hourly`, and thus try to use that pointer to call the function it points to:

```
Vtable_Hourly[0]((struct HourlyEmployee *)&h);
```

This makes sense in our head: the element 0 of `Vtable_Hourly` is a pointer to `Speak_Hourly`, which wants a single parameter of type `struct Employee *`. Unfortunately C sees `Vtable_Hourly[0]` as a pointer to `void`, not to a function; we have to tell the compiler that it actually points to a function that returns `void` and takes a single parameter of type `struct Employee *`. A cast will accomplish this task:

```
((void (*)(struct Employee*))Vtable_Hourly[0])((struct Employee *)&h);
```

which tells the compiler to invoke the function pointed to by `Vtable_Hourly[0]`, passing it the address of `h` as its parameter. Success!

All we need to do now is add a new member variable to our structs: a `void**` pointer to a vtable appropriate to the class. Create one vtable variable globally for all derived `Employee` types, point the vtable pointers to the appropriate global tables in constructor methods, and voila! We can now call a virtual function through a base class pointer by accessing the appropriate index from the vtable associated with the variable and invoking that function by hand.

## Program

Implement the following code:

1. Create `Employee`, `HourlyEmployee`, and `CommissionEmployee` structs as described above. The first member of each struct should be a `void** vtable` variable.
2. Implement the `Speak_Hourly` and `Speak_Commission` functions for the `HourlyEmployee` and `CommissionEmployee` structs as global functions.
3. Add `HourlyEmployee` and `CommissionEmployee` implementations of a `GetPay` function, which takes an `Employee` pointer and returns a `double` as such:
  - (a) The pay for an hourly employee is the number of hours multiplied by their hourly rate.
  - (b) The pay for a commission employee is 10% of their total sales, plus 40,000.
4. Create a `Vtable_XX` global array for each employee derived type, initialized with pointers to the appropriate `Speak` and `GetPay` functions for that type, **in that order**.
5. Implement `Construct_XX` functions for `HourlyEmployee` and `CommissionEmployee`, which initialize the member variables to 0 values, and **most importantly**, sets the employee's vtable pointer to the appropriate global `Vtable_XX` variable.
6. Add a new "class" `SeniorSalesman`, which "derives" from `CommissionEmployee` by duplicating all of `CommissionEmployee`'s member variables. `SeniorSalesman` will override the `GetPay` method but will use `CommissionEmployee`'s version of `Speak`. Create a vtable and constructor for the `SeniorSalesman` class, and a `GetPay` method that returns 20% of the salesman's sales, plus 50,000, plus another 5% of sales if the employee is at least 40 years old.

Then write a main program that does the following:

1. Declare an `Employee` pointer.
2. Ask the user to choose either an hourly employee, a commission employee, or a senior salesman.
3. Use `malloc` to create space for the appropriate employee (since this is a dynamic memory need).
  - (a) Ask the user how old the employee is.

- (b) If the user selects an hourly, ask them for the employee's pay rate and hours.
  - (c) If the user selects a commission or a senior salesman, ask for the employee's amount of sales.
  - (d) Use your `Construct_` function to initialize the memory given back from `malloc`, passing the appropriate parameters for the type selected.
  - (e) Point your `Employee` to the variable.
  - (f) **Past this point of the program, you can have no code referring to `CommissionEmployee`, `HourlyEmployee`, or `SeniorSalesman` explicitly – everything must be through `Employee` pointers.**
- 4. Tell the `Employee` to speak, then inform the user how much money they make.
    - (a) To do this, you will access the `vtable` pointer from the `Employee` pointer in your main, index it to the appropriate position for each function, cast that pointer as described above, and invoke the function it is pointing to.
  - 5. Use `free` to free the variable you created.

## CECS 424H Requirements

Nothing ridiculous; you're simply going to take your Lab 2 custom allocator and use `my_alloc` and `my_free` in the place of `malloc` and `free` :).

## Deliverables

Turn in the following when the lab is due:

- 1. A printed copy of your code, **printed from your IDE when possible**. If you cannot print from your editor, copy your code into Notepad or another program with a fixed-width (monospace) font and print from there.
- 2. A printed copy of the output of your program, where you choose:
  - (a) an hourly employee 25 years old making \$9.50 an hour working 90 hours.
  - (b) a commission employee 30 years old with \$80,000 in sales.
  - (c) a senior salesman 50 years old with \$100,000 in sales.